



Experiment No.4
Experiment on Hadoop Map-Reduce
Date of Performance: 07/08/23
Date of Submission: 14/08/23

AIM: -To write a program to implement a word count program using MapReduce.

THEORY:

WordCount is a simple program which counts the number of occurrences of each word in a given text input data set. WordCount fits very well with the MapReduce programming model making it a great example to understand the Hadoop Map/Reduce programming style. The implementation consists of three main parts:

1. Mapper
2. Reducer
3. Driver



Step-1. Write a Mapper

A Mapper overrides the `map()` function from the Class `"org.apache.hadoop.mapreduce.Mapper"` which provides `<key, value>` pairs as the input. A Mapper implementation may output `<key,value>` pairs using the provided Context .

Input value of the WordCount Map task will be a line of text from the input data file and the key would be the line number `<line_number, line_of_text>` . Map task outputs `<word, one>` for each word in the line of text. Pseudo-code void Map (key, value){ for each word x in value:

```
output.collect(x,1);  
}
```

Step-2. Write a Reducer

A Reducer collects the intermediate `<key,value>` output from multiple map tasks and assemble a single result. Here, the WordCount program will sum up the occurrence of each word to pairs as `<word, occurrence>`. Pseudo-code

```
void Reduce (keyword, <list of value>){ for each  
x in <list of value>:  
sum+=x;  
final_output.collect(keyword, sum);  
}
```

Code:

```
import java.io.IOException; import  
java.util.StringTokenizer; import  
org.apache.hadoop.io.IntWritable; import  
org.apache.hadoop.io.LongWritable; import  
org.apache.hadoop.io.Text; import  
org.apache.hadoop.mapreduce.Mapper; import  
org.apache.hadoop.mapreduce.Reducer; import  
org.apache.hadoop.conf.Configuration; import  
org.apache.hadoop.mapreduce.Job;
```



```
import org.apache.hadoop.mapreduce.lib.input.TextInputFormat;
import org.apache.hadoop.mapreduce.lib.output.TextOutputFormat;
import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;
import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;
import org.apache.hadoop.fs.Path; public class WordCount
{
    public static class Map extends Mapper<LongWritable,Text,Text,IntWritable> { public
    void map(LongWritable key, Text value,Context context) throws
    IOException,InterruptedException{
        String line = value.toString();
        StringTokenizer tokenizer = new StringTokenizer(line);
        while (tokenizer.hasMoreTokens()) {
            value.set(tokenizer.nextToken()); context.write(value,
            new IntWritable(1));
        }
    }

    public static class Reduce extends Reducer<Text,IntWritable,Text,IntWritable> {
        public void reduce(Text key, Iterable<IntWritable> values,Context context)
        throws IOException,InterruptedException { int sum=0; for(IntWritable x: values)
        { sum+=x.get();
        }
        context.write(key, new IntWritable(sum));
    }
}

public static void main(String[] args) throws Exception
{ Configuration conf= new Configuration(); Job job =
new Job(conf,"My Word Count Program");
```



```
job.setJarByClass(WordCount.class);
job.setMapperClass(Map.class);
job.setReducerClass(Reduce.class);
job.setOutputKeyClass(Text.class);
job.setOutputValueClass(IntWritable.class);
job.setInputFormatClass(TextInputFormat.class);
job.setOutputFormatClass(TextOutputFormat.class);

Path outputPath = new Path(args[1]);

//Configuring the input/output path from the filesystem into the job

FileInputFormat.addInputPath(job, new Path(args[0])); FileOutputFormat.setOutputPath(job,
new Path(args[1]));

//deleting the output path automatically from hdfs so that we don't have to delete
it explicitly

outputPath.getFileSystem(conf).delete(outputPath); //exiting
the job only if the flag value becomes false

System.exit(job.waitForCompletion(true) ? 0 : 1);
}
}
```

OUTPUT:



Vidyavardhini's College of Engineering & Technology

Department of Computer Engineering

BDA Lab Manual 23-24

Namenode information

localhost:9870/dfshealth.html#tab-overview

Hadoop

Overview

Datanodes

Datanode Volume Failures

Snapshot

Startup Progress

Utilities

Overview 'localhost:9820' (active)

Started:	Wed Sep 13 04:30:53 +0530 2023
Version:	3.2.4, r7e5d9983b388e372fe640f21f048f2f2ae6e9eba
Compiled:	Tue Jul 12 17:28:00 +0530 2022 by ubuntu from branch-3.2.4
Cluster ID:	CID-146566e0-df7a-44ee-a644-d41c94627871
Block Pool ID:	BP-1532262397-192.168.12.89-1692767105768

Summary

Security is off.

Safemode is off.

3 files and directories, 1 blocks (1 replicated blocks, 0 erasure coded block groups) = 4 total filesystem object(s).

Heap Memory used 93.19 MB of 204.5 MB Heap Memory. Max Heap Memory is 889 MB.

Non Heap Memory used 51.98 MB of 53.3 MB Committed Non Heap Memory. Max Non Heap Memory is <unbounded>.

Configured Capacity:	417.65 GB
Configured Remote Capacity:	0 B
DFS Used:	345 B (0%)

File Edit View Navigate Code Refactor Build Run Tools VCS Window Help

WordCount - pom.xml (WordCount)

Project

WordCount

WordCount - pom.xml (WordCount)

1 <?xml version="1.0" encoding="UTF-8"?>

2 <project xmlns="http://maven.apache.org/POM/4.0.0"

3 xmlns:xsi="http://www.w3.org/2001/XMLSchema-instance"

4 xsi:schemaLocation="http://maven.apache.org/POM/4.0.0 http://maven.apache.org/xsd/maven-4.0.0.xsd">

5 <modelVersion>4.0.0</modelVersion>

6

7 <groupId>org.samarth</groupId>

8 <artifactId>WordCount</artifactId>

9 <version>1.0-SNAPSHOT</version>

10

11 <properties>

12 <maven.compiler.source>11</maven.compiler.source>

13 <maven.compiler.target>11</maven.compiler.target>

14 <project.build.sourceEncoding>UTF-8</project.build.sourceEncoding>

15 </properties>

16

17 <dependencies>

18 <dependency>

19 <groupId>org.apache.hadoop</groupId>

20 <artifactId>hadoop-common</artifactId>

21 <version>3.3.3</version>

22 </dependency>

23 <dependency>

24 <groupId>org.apache.hadoop</groupId>

25 <artifactId>hadoop-mapreduce-client-core</artifactId>

26 <version>3.3.3</version>

27 </dependency>

28 </dependencies>

29

30 </project>

project > dependencies > dependency

Version Control TODO Problems Terminal Services Build Dependencies

Suggested plugin Protocol Buffers available for dependency 'java:com.google.protobuf:protobuf-java'.// Configure plugins...// Don't suggest this plugin (moments ago)

27:22 LF UTF-8 4 spaces



Vidyavardhini's College of Engineering & Technology

Department of Computer Engineering

The image displays two screenshots of the IntelliJ IDEA IDE, showing the implementation of a WordCount application using MapReduce.

Top Screenshot: WC_Mapper.java

```
1 package org.samarth;
2 import java.io.IOException;
3 import java.util.StringTokenizer;
4 import org.apache.hadoop.io.IntWritable;
5 import org.apache.hadoop.io.LongWritable;
6 import org.apache.hadoop.io.Text;
7 import org.apache.hadoop.mapred.MapReduceBase;
8 import org.apache.hadoop.mapred.Mapper;
9 import org.apache.hadoop.mapred.OutputCollector;
10 import org.apache.hadoop.mapred.Reporter;
11
12 public class WC_Mapper extends MapReduceBase implements Mapper<LongWritable,Text,Text,IntWritable> {
13     private final static IntWritable one = new IntWritable(1);
14     private Text word = new Text();
15
16     public void map(LongWritable key, Text value, OutputCollector<Text, IntWritable> output,
17                     Reporter reporter) throws IOException {
18         String line = value.toString();
19         StringTokenizer tokenizer = new StringTokenizer(line);
20         while (tokenizer.hasMoreTokens()) {
21             word.set(tokenizer.nextToken());
22             output.collect(word, one);
23         }
24     }
25 }
```

Bottom Screenshot: WC_Reducer.java

```
1 package org.samarth;
2 import java.io.IOException;
3 import java.util.Iterator;
4 import org.apache.hadoop.io.IntWritable;
5 import org.apache.hadoop.io.Text;
6 import org.apache.hadoop.mapred.MapReduceBase;
7 import org.apache.hadoop.mapred.OutputCollector;
8 import org.apache.hadoop.mapred.Reducer;
9 import org.apache.hadoop.mapred.Reporter;
10
11 public class WC_Reducer extends MapReduceBase implements Reducer<Text,IntWritable,Text,IntWritable> {
12     public void reduce(Text key, Iterator<IntWritable> values, OutputCollector<Text,IntWritable> output,
13                       Reporter reporter) throws IOException {
14         int sum=0;
15         while (values.hasNext()) {
16             sum+=values.next().get();
17         }
18         output.collect(key, new IntWritable(sum));
19     }
20 }
```



Vidyavardhini's College of Engineering & Technology

Department of Computer Engineering

The image shows a screenshot of an IDE (IntelliJ IDEA) and a Windows Command Prompt window. The IDE window displays the code for a MapReduce application named 'WordCount'. The code is written in Java and uses the Hadoop MapReduce framework. The code is as follows:

```
import java.io.IOException;
import org.apache.hadoop.fs.Path;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapred.FileInputFormat;
import org.apache.hadoop.mapred.FileOutputFormat;
import org.apache.hadoop.mapred.JobClient;
import org.apache.hadoop.mapred.JobConf;
import org.apache.hadoop.mapred.TextInputFormat;
import org.apache.hadoop.mapred.TextOutputFormat;

public class WC_Runner {
    public static void main(String[] args) throws IOException{
        JobConf conf = new JobConf(WC_Runner.class);
        conf.setJobName("WordCount");
        conf.setOutputKeyClass(Text.class);
        conf.setOutputValueClass(IntWritable.class);
        conf.setMapperClass(WC_Mapper.class);
        conf.setCombinerClass(WC_Reducer.class);
        conf.setReducerClass(WC_Reducer.class);
        conf.setInputFormat(TextInputFormat.class);
        conf.setOutputFormat(TextOutputFormat.class);
        FileInputFormat.setInputPaths(conf, new Path(args[0]));
        FileOutputFormat.setOutputPath(conf, new Path(args[1]));
        JobClient.runJob(conf);
    }
}
```

The Command Prompt window shows the following commands being executed:

```
Microsoft Windows [Version 10.0.22000.2295]
(c) Microsoft Corporation. All rights reserved.

C:\Users\admin>cd Desktop

C:\Users\admin\Desktop>hadoop fs -mkdir /input

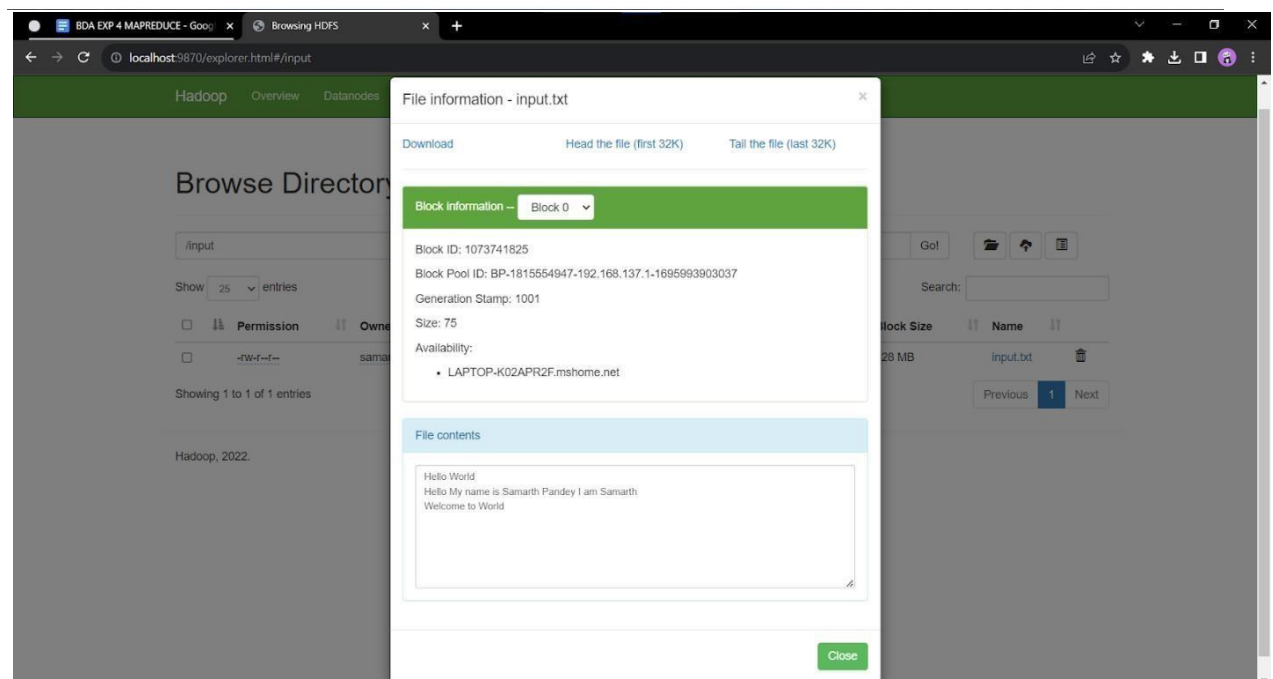

C:\Users\admin\Desktop>hadoop fs -put input.txt /input

C:\Users\admin\Desktop>
```



Vidyavardhini's College of Engineering & Technology

Department of Computer Engineering





The screenshot shows a Windows Command Prompt window and a web browser window. The Command Prompt displays the execution of a Hadoop job, including file creation, job submission, and completion status. The web browser shows the Hadoop file system interface, displaying a directory listing of files named 'input', 'output', and 'tmp'.

```
Command Prompt
C:\Users\samar\Desktop\hadoop fs -mkdir /input
C:\Users\samar\Desktop\hadoop fs -put input.txt /input
C:\Users\samar\Desktop\hadoop jar C:\Users\samar\IdeaProjects\WordCount\target\hadoop-mapreduce-3.2.4.jar wordcount /input/input.txt /output
2023-09-29 18:57:08,319 INFO client.RMProxy: Connecting to ResourceManager at 0.0.0.0:8032
2023-09-29 18:57:09,763 INFO mapreduce.JobResourceUploader: Disabling Erasure Coding for path: /tmp/hadoop-yarn/staging/samar/.staging/job_1695993949979_0001
2023-09-29 18:57:10,607 INFO mapreduce.JobSubmitter: number of splits:1
2023-09-29 18:57:11,007 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1695993949979_0001
2023-09-29 18:57:11,010 INFO mapreduce.JobSubmitter: Executing with tokens: []
2023-09-29 18:57:11,299 INFO conf.Configuration: resource-types.xml not found
2023-09-29 18:57:11,300 INFO resource.ResourceUtils: Unable to find 'resource-types.xml'.
2023-09-29 18:57:11,723 INFO Impl.YarnClientImpl: Submitted application application_1695993949979_0001
2023-09-29 18:57:11,814 INFO mapreduce.Job: The url to track the job: http://AAPIDP-K02APR2F:8088/proxy/application_1695993949979_0001/
2023-09-29 18:57:11,816 INFO mapreduce.Job: Running job: job_1695993949979_0001
2023-09-29 18:57:27,135 INFO mapreduce.Job: Job job_1695993949979_0001 running in uber mode : false
2023-09-29 18:57:27,136 INFO mapreduce.Job: map 0% reduce 0%
2023-09-29 18:57:35,308 INFO mapreduce.Job: map 100% reduce 0%
2023-09-29 18:57:42,413 INFO mapreduce.Job: map 100% reduce 100%
2023-09-29 18:57:44,434 INFO mapreduce.Job: Job job_1695993949979_0001 completed successfully
2023-09-29 18:57:45,177 INFO mapreduce.Job: Counters: 54
File System Counters
  FILE: Number of bytes read=126
  FILE: Number of bytes written=478809
  FILE: Number of read operations=0
  FILE: Number of large read operations=0
  FILE: Number of write operations=0
  HDFS: Number of bytes read=177
  HDFS: Number of bytes written=76
  HDFS: Number of read operations=8
  HDFS: Number of large read operations=0
  HDFS: Number of write operations=2
  HDFS: Number of bytes read erasure-coded=0
Job Counters
  Launched map tasks=1
  Launched reduce tasks=1
  Data-local map tasks=1
  Total time spent by all maps in occupied slots (ms)=5488
  Total time spent by all reduces in occupied slots (ms)=5838
  Total time spent by all map tasks (ms)=5488
  Total time spent by all reduce tasks (ms)=5838
  Total vcore-milliseconds taken by all map tasks=5488
  Total vcore-milliseconds taken by all reduce tasks=5838
  Total megabyte-milliseconds taken by all map tasks=5619712
  Total megabyte-milliseconds taken by all reduce tasks=5978112
Map-Reduce Framework
  Map input records=3
```

The web browser shows the Hadoop file system interface, displaying a directory listing of files named 'input', 'output', and 'tmp'.

Browse Directory

/

Go!

Show

25

entries

Search:

<input type="checkbox"/>		Permission		Owner		Group		Size		Last Modified		Replication		Block Size		Name	
<input type="checkbox"/>		drwxr-xr-x		samar		supergroup		0 B		Sep 29 18:56		0		0 B		input	
<input type="checkbox"/>		drwxr-xr-x		samar		supergroup		0 B		Sep 29 18:57		0		0 B		output	
<input type="checkbox"/>		drwxr-xr-x		samar		supergroup		0 B		Sep 29 18:57		0		0 B		tmp	

Showing 1 to 3 of 3 entries

Previous

1

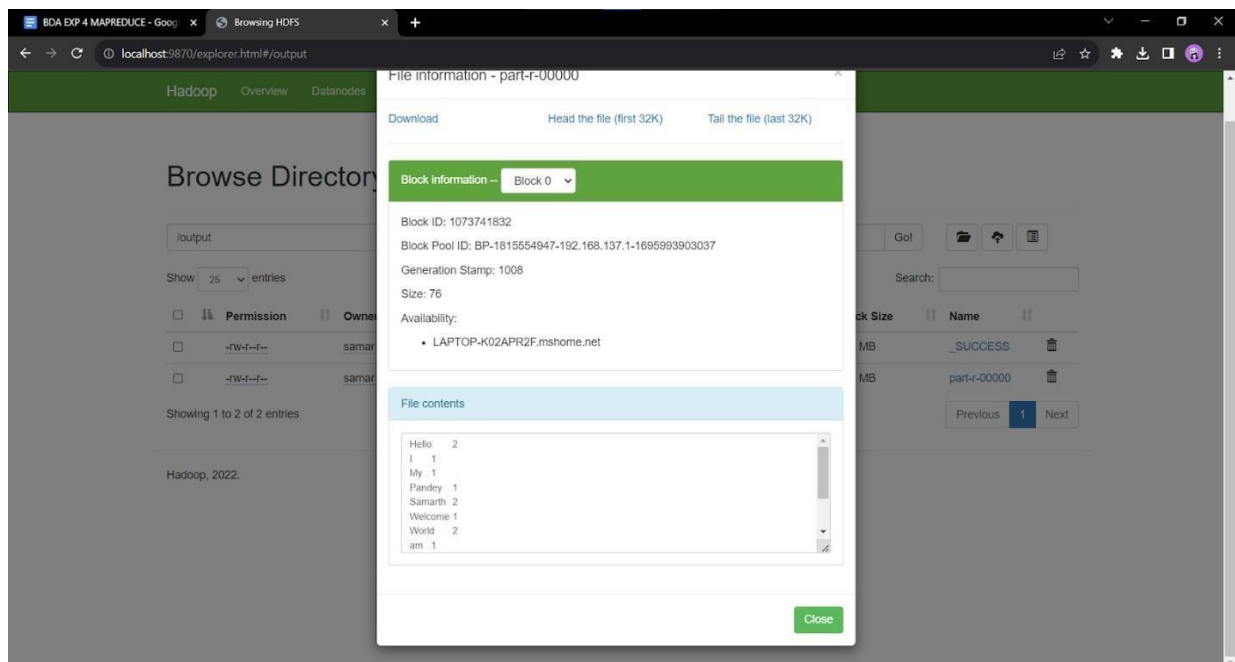
Next

Hadoop, 2022.



Vidyavardhini's College of Engineering & Technology

Department of Computer Engineering



CONCLUSION :

The attempt to use MapReduce to construct a word count programme was a resounding success. When dealing with large datasets, MapReduce clearly showed its scalability and efficiency, assigning tasks to different nodes effectively to enable parallel processing. In particular, it demonstrated how it might ensure the integrity of data processing in dispersed systems. A wide range of developers can use MapReduce thanks to its simple design and simple-to-understand mapper and reducer functionalities. This experiment has real-world applications for more complex data processing tasks like log analysis and machine learning. The program's efficiency can also be increased by incorporating performance optimisation techniques like combiners and partitioners.